


VAN EMDE-BOAS SEARCH TREES

Dominik Bez

- Implement **van Emde-Boas tree** as a fast `std::set` alternative
- Operations: insert, remove, locate/lower_bound in $O(\log \log n)$
 - **locate**(x) = $\min\{y \in VEB \mid y \geq x\}$
- Supporting unsigned/signed integers and floating point values
- Different variations
 - **Generic** van Emde-Boas tree (VEB)
 - Efficient specialized **32 bit** implementation (VEB32)
 - Various **concurrent** versions of VEB32




- Recursion ends for $K \leq 6 \Rightarrow$ store a single 64 bit integer
 - Use bitwise and least/most significant bit operations
- Recursion also ends for a **single value** in bottom data structure
 - Least significant bit in pointer denotes this case
 - $K < \text{sizeof}(\text{void}^*) \Rightarrow$ store data **in the pointer itself**
 - Otherwise, store the element **directly on the heap**

GENERIC VAN EMDE-BOAS TREE - RECURSION BASE CASE


```
1 uint64_t storage;
2
3 void insert(const uint32_t value) noexcept {
4     storage |= 1ULL << value;
5 }
6
7 void remove(const uint32_t value) noexcept {
8     storage &= ~(1ULL << value);
9 }
10
11 uint32_t locate(const uint32_t value) const noexcept {
12     return std::countr_zero(storage & ~((1ULL << value) - 1));
13 }
```

GENERIC VAN EMDE-BOAS TREE - EVALUATION

Time of 10000 Lookups in a Tree with 'Size' Elements (uniform distribution)




32 BIT VAN EMDE-BOAS TREE



- Exploiting 64 bit processors
 - Top tree manages upper 18 bits instead of 16
 - Bottom trees manage lower 14 bits
 - Their top tree is flat and manages 8 bits
 - The remaining 6 bits are managed by a 64 bit integer again
 - \Rightarrow One level less than in the original paper
- Single values are stored directly in the pointer again
- Consumes at most 1 GB of memory compared to up to 100 GB

32 BIT VAN EMDE-BOAS TREE - EVALUATION

Time of 10000 Lookups in a Tree with 'Size' Elements (uniform distribution)



- Locked top
 - Top data structure **completely locked** (shared lock for locate)
 - Once the correct bottom data structure is found, it is locked instead
- Locked fine-grained
 - One mutex per 64 bottom data structures (one integer in top data structure)
- „Lockless“
 - Top data structure is **lockless**
 - **Atomically swap** pointer to bottom data structure with reserved value to reserve it (wait with spinlock)
- Remove always uniquely locks a **shared mutex**

CONCURRENT IMPLEMENTATIONS - EVALUATION

Speedup over uVEB32 to Insert 'Size' Elements Parallel (uniform distribution; 2097152 Elements)

